

Logit

Suppose that agents can choose an action $y \in \{0, 1\}$. The payoffs to each are

$$\begin{aligned}u_0(X, \epsilon) &= \epsilon_{i0} \\ u_1(X, \epsilon) &= x_i^\top \beta + \epsilon_{i1}\end{aligned}$$

The shocks $\epsilon_{i0}, \epsilon_{i1}$ are distributed as iid Type-I extreme value with mean 0 and scale parameter 1. Agents choose $y_i = 1$ if $u_1 \geq u_0$ and y_0 if $u_1 < u_0$

Denote $\text{cdf}(z) = F(z)$. Let choice $y_i \in \{0, 1\}$.

Likelihood is

$$\log L(y|X) = \sum_i \log F(2(y_i - 1)x_i^\top \beta)$$

Score is

$$\nabla_\beta \log L(y_i|x_i) = [y_i - F(x_i^\top \beta)] x_i$$

Information matrix is

$$\left[\sum_i \nabla \log L_i \nabla \log L_i^\top \right] \rightarrow \text{Var}(\beta)$$

We can use the following functions from **StatsFuns.jl**

```
logsumexp    # log(exp(x) + exp(y)) or log(sum(exp(x)))
softmax      # exp(x_i) / sum(exp(x)), for i
```

```
• md"
• # Logit
•
• Suppose that agents can choose an action $y \in \{0,1\}$. The payoffs to each are
•
• ```math
• \begin{align*}
• u_0(X, \epsilon) &= \epsilon_{i0} \\
• u_1(X, \epsilon) &= x_i^\top \beta + \epsilon_{i1}
• \end{align*}
• ```
•
•
```

- The shocks $\epsilon_{i0}, \epsilon_{i1}$ are distributed as iid Type-I extreme value with mean 0 and scale parameter 1. Agents choose $y_i=1$ if $u_1 \geq u_0$ and $y_i=0$ if $u_1 < u_0$
- Denote $F(z) = \Pr(y=1|X=z)$. Let choice $y_i \in \{0,1\}$.
- Likelihood is
 - $$\log L(y|X) = \sum_i \log F(y_i x_i^{\top \beta})$$
- Score is
 - $$\nabla_{\beta} \log L(y_i|x_i) = [y_i - F(x_i^{\top \beta})] x_i$$
- Information matrix is
 - $$-\mathbb{E} \left[\sum_i \nabla_{\beta} \log L_i \nabla_{\beta} \log L_i^{\top} \right] \rightarrow \text{Var}(\beta)$$
- We can use the following functions from [`StatsFuns.jl`] (<https://github.com/JuliaStats/StatsFuns.jl>)


```

julia
logsumexp      # log(exp(x) + exp(y)) or log(sum(exp(x)))
softmax        # exp(x_i) / sum(exp(x)), for i
      
```

- begin
 - using Random: seed!
 - using StatsFuns: logsumexp, softmax
 - using LinearAlgebra: diag
 - using Optim
 - using Distributions
 - using DataFrames
 - using StatsBase: countmap
 - *# autodiff instead of finite diff?*
 - using FiniteDiff: finite_difference_gradient
- end

MersenneTwister(1234)

- `seed!(1234)`

```

1000x2 Matrix{Float64}:
0.0  -2.89099
0.0   1.2706
0.0  -2.74435
0.0  -2.06214
0.0  -0.363533
0.0   1.73134
0.0  -0.725822
⋮
0.0  -4.08263
0.0   0.8717
0.0  -5.16962
      
```

```
0.0 -1.80468
0.0 -0.734464
0.0 3.2252
```

```
• begin
•     nob = 1_000
•      $\beta$  = [1.0, -2.0, 1.0, 0.5]
•     k = length( $\beta$ )
•     X = randn(nob, k);
•
•     # choice utilities
•     u0 = zeros(nob)
•     u1 = X* $\beta$ 
•     u = hcat(u0, u1)
• end
```

```
prob_actions = 1000x2 Matrix{Float64}:
      0.947399  0.0526007
      0.219155  0.780845
      0.939594  0.0604063
      0.887168  0.112832
      0.589895  0.410105
      0.150417  0.849583
      0.673888  0.326112
      ⋮
      0.983417  0.0165835
      0.294901  0.705099
      0.994345  0.00565458
      0.858717  0.141283
      0.675784  0.324216
      0.0382285 0.961771
```

```
• # multinomial logit probabilities
• prob_actions = mapslices(softmax, u; dims=2)
```

```
cum_prob = 1000x2 Matrix{Float64}:
      0.947399  1.0
      0.219155  1.0
      0.939594  1.0
      0.887168  1.0
      0.589895  1.0
      0.150417  1.0
      0.673888  1.0
      ⋮
      0.983417  1.0
      0.294901  1.0
      0.994345  1.0
      0.858717  1.0
      0.675784  1.0
      0.0382285 1.0
```

```
• cum_prob = cumsum(prob_actions; dims=2)
```

```
• @assert all(cum_prob[:,2] .≈ 1)
```

```
y =
```

```
Int64[0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, more ,
```

```
• # instead of simulating random type-1 extreme values, we just
• # use a uniform variable and the CDF
```

```
• y = [searchsortedfirst(row, rand()) for row in eachrow(cum_prob) ] .-1
```

```
Dict{0 ⇒ 509, 1 ⇒ 491}
```

```
• countmap(y)
```

checksizes (generic function with 1 method)

```
• function checksizes(y, X, theta)
•     @assert eltype(y) <: Integer
•     n,k = size(X)
•     n == length(y) || throw(DimensionMismatch())
•     k == length(theta) || throw(DimensionMismatch())
•     return n, k
• end
```

loglik (generic function with 1 method)

```
• function loglik(y, X, theta)
•     n,k = checksizes(y,X,theta)
•     ff(z) = logcdf(Logistic(), z)
•
•     # see footnote 6 on p. 778 in Greene 6th ed for this shortcut
•     q = 2 .* y .- 1
•     u1 = X*theta
•     LL = sum(ff.(q.*u1))
•
•     return -LL # I *think* you'll need to flip sign to maximize
• end
```

dloglik! (generic function with 1 method)

```
• function dloglik!(grad, y, X, theta)
•
•     n,k = checksizes(y,X,theta)
•     k == length(grad) || throw(DimensionMismatch())
•
•     u1 = X*theta
•     ff(z) = cdf(Logistic(), z)
•     g = y .- ff.(u1) # as per Greene 6th ed p. 779
•
•     grad .= -vec(sum(g .* X; dims=1))
•     return grad
• end
```

dloglik (generic function with 1 method)

```
• dloglik(y, X, theta) = dloglik!(similar(theta), y, X, theta)
```

informationmatrix (generic function with 1 method)

```
• function informationmatrix(y, X, theta)
•
•     n,k = checksizes(y,X,theta)
•     infomatrix = zeros(k,k)
•
•     u1 = X*theta
•     ff(z) = cdf(Logistic(), z)
•     g = y .- ff.(u1) # as per Greene 6th ed p. 779
•
•     infomatrix = (g .* X)' * (g .* X)
```

```

    return infomatrix # maybe flip signs?
end

```

g! (generic function with 1 method)

```

• # closures wrap likelihood & gradient
• begin
•   f(thet) = loglik(y,X,thet)
•   g!(grad,thet) = dloglik!(grad,y,X,thet)
• end

```

```
theta0 = Float64[0.0, 0.0, 0.0, 0.0]
```

```

• # initial guess
• theta0 = zeros(k)

```

```
Float64[1.46633e-9, -1.11e-8, 6.52823e-9, -1.92486e-9]
```

```

• # Check gradient against autodiff
• begin
•   fdgrad = finite_difference_gradient(f, theta0, Val{:central})
•   @assert fdgrad ≈ dloglik(y,X,theta0)
•   fdgrad .- dloglik(y,X,theta0)
• end

```

```
res = * Status: success
```

```

* Candidate solution
  Final objective value:      3.917137e+02

```

```

* Found with
  Algorithm:      BFGS

```

```

* Convergence measures
  |x - x'|          = 5.80e-10 ≠ 0.0e+00
  |x - x'|/|x'|     = 2.63e-10 ≠ 0.0e+00
  |f(x) - f(x')|     = 5.68e-14 ≠ 0.0e+00
  |f(x) - f(x')|/|f(x')| = 1.45e-16 ≠ 0.0e+00
  |g(x)|            = 1.36e-12 ≤ 1.0e-08

```

```

* Work counters
  Seconds run:      2 (vs limit Inf)
  Iterations:       14
  f(x) calls:       42
  ∇f(x) calls:      42

```

```
• res = optimize(f, g!, theta0, BFGS(), Optim.Options(;show_trace=true))
```

	beta	betahat	tstat	se	pval
1	1.0	0.991798	10.2166	0.0970772	1.67126e-24
2	-2.0	-2.20044	-15.7587	0.139634	5.9898e-56
3	1.0	0.990966	9.66605	0.10252	4.20272e-22

	beta	betahat	tstat	se	pval
4	0.5	0.510733	5.34361	0.0955783	9.1112e-8

```

• begin
•   theta1 = res.minimizer # should be about  $\beta$ 
•   vcov = informationmatrix(y, X, theta1)
•   vcovinv = inv(vcov)
•   stderr = sqrt(diag(vcovinv))
•   tstats = theta1 ./ stderr
•   pvals = map(z -> 2 .* cdf(Normal(), -abs(z)), tstats)
•
•   DataFrame(beta =  $\beta$ , betahat = theta1, tstat=tstats, se = stderr, pval=pvals)
• end

```