

Logit

Suppose that agents can choose an action $y \in \{0, 1\}$. The payoffs to each are

$$\begin{aligned} u_0(X, \epsilon) &= \epsilon_{i0} \\ u_1(X, \epsilon) &= x_i^\top \beta + \epsilon_{i1} \end{aligned}$$

The shocks $\epsilon_{i0}, \epsilon_{i1}$ are distributed as iid Type-I extreme value with mean 0 and scale parameter 1. Agents choose $y_i = 1$ if $u_1 \geq u_0$ and y_0 if $u_1 < u_0$

```
• md"
• # Logit
•
• Suppose that agents can choose an action $y \in \{0,1\}$. The payoffs to each are
•
• ```math
• \begin{align*}
• u_0(X, \epsilon) &= \epsilon_{i0} \\
• u_1(X, \epsilon) &= x_i^\top \beta + \epsilon_{i1}
• \end{align*}
• ```
•
• The shocks $\epsilon_{i0}, \epsilon_{i1}$ are distributed as iid Type-I extreme value
  with mean 0 and scale parameter 1. Agents choose $y_i=1$ if $u_1 \geq u_0$ and $y_0$
  if $u_1 < u_0$
• "
```

Denote $\text{cdf}(z) = F(z)$. Let choice $y_i \in \{0, 1\}$. For a symmetric distribution and a binary discrete choice model, we can use this shortcut (trick is in Greene's Econometrics tome, Greene 6th ed p. 779):

$$\log L(y|X) = \sum_i \log F(2(y_i - 1)x_i^\top \beta)$$

Score is a vector

$$\nabla_\beta \log L(y_i|x_i) = [y_i - F(x_i^\top \beta)] x_i$$

Information matrix is

$$\left[\sum_i \nabla \log L_i \nabla \log L_i^\top \right] \rightarrow \text{Var}(\beta)$$

```
• md"
• Denote $cdf(z) = F(z)$. Let choice $y_i \in \{0,1\}$. For a symmetric distribution
  and a binary discrete choice model, we can use this shortcut (trick is in Greene's
  Econometrics tome, Greene 6th ed p. 779):
• ```math
```

- $\log L(y|X) = \sum_i \log F(\text{2}(y_{i-1}) x_i^{\text{top}} \beta)$
- $\log L(y|X) = \sum_i \log F(\text{2}(y_{i-1}) x_i^{\text{top}} \beta)$
- Score is a vector
- $\nabla_{\beta} \log L(y_i|x_i) = [y_i - F(x_i^{\text{top}} \beta)] x_i$
- Information matrix is
- $\left[\sum_i \nabla \log L_i \nabla \log L_i^{\text{top}} \right] \rightarrow \text{Var}(\beta)$
- "

We can vectorize stuff to make it simpler. The . means element-by-element operations a la MATLAB/Julia. Define

$$\mathbf{q} \equiv \mathbf{2} \cdot \mathbf{y} - \mathbf{1}$$

Then

$$\log L(y|X) = \sum_i \log F(\mathbf{q} \cdot X \beta)$$

Matrix of scores

$$\frac{\partial \log L_i}{\partial \beta} = (y_i - F(X \beta)) \cdot X$$

Information matrix

$$\left(\frac{\partial \log L_i}{\partial \beta} \right)^{\top} \frac{\partial \log L_i}{\partial \beta} \rightarrow \text{Var}(\beta)$$

- md"
- We can vectorize stuff to make it simpler. The \$. \$ means element-by-element operations a la MATLAB/Julia. Define
- $\mathbf{q} \equiv \mathbf{2} \cdot \mathbf{y} - \mathbf{1}$
- Then
- $\log L(y|X) = \sum_i \log F(\mathbf{q} \cdot X \beta)$
- Matrix of scores
- $\frac{\partial \log L_i}{\partial \beta} = (y_i - F(X \beta)) \cdot X$
- Information matrix
- $\left(\frac{\partial \log L_i}{\partial \beta} \right)^{\top} \frac{\partial \log L_i}{\partial \beta} \rightarrow \text{Var}(\beta)$
- "

In the version below, we use the `Distributions.jl` package, which means we could actually change to a binary probit just by swapping out the distribution from `Logistic` to `Normal`.

Alternatively, for lower-level control, we can use the following functions from `StatsFuns.jl`. This is useful for more computationally intensive work with multinomial discrete choice.

```
logsumexp    # log(exp(x) + exp(y)) or log(sum(exp(x)))
softmax      # exp(x_i) / sum(exp(x)), for i
```

```
• md"
• In the version below, we use the `Distributions.jl` package, which means we could
  actually change to a binary probit just by swapping out the distribution from
  `Logistic` to `Normal`.
•
• Alternatively, for lower-level control, we can use the following functions from
  [`StatsFuns.jl`](https://github.com/JuliaStats/StatsFuns.jl). This is useful for more
  computationally intensive work with multinomial discrete choice.
•
• ```julia
• logsumexp    # log(exp(x) + exp(y)) or log(sum(exp(x)))
• softmax      # exp(x_i) / sum(exp(x)), for i
• ```
• "
```

```
• begin
•     # import the entire package
•     using Optim
•     using Distributions
•     using DataFrames
•
•     # import just a few functions
•     using Random: seed!
•     using StatsFuns: logsumexp, softmax
•     using LinearAlgebra: diag
•     using StatsBase: countmap
•
•     # autodiff instead of finite diff?
•     using FiniteDiff: finite_difference_gradient
• end
```

MersenneTwister(1234)

```
• # set seed for random # generator
• seed!(1234)
```

1000×2 Matrix{Float64}:

```
0.0 -2.89099
0.0  1.2706
0.0 -2.74435
0.0 -2.06214
0.0 -0.363533
0.0  1.73134
0.0 -0.725822
⋮
0.0 -4.08263
0.0  0.8717
```

```
0.0 -5.16962
0.0 -1.80468
0.0 -0.734464
0.0 3.2252
```

```
• begin
•     nob = 1_000
•      $\beta$  = [1.0, -2.0, 1.0, 0.5]
•     k = length( $\beta$ )
•     X = randn(nob, k);
•
•     # choice utilities
•     u0 = zeros(nob)
•     u1 = X* $\beta$ 
•     u = hcat(u0, u1)
• end
```

```
prob_actions = 1000x2 Matrix{Float64}:
      0.947399  0.0526007
      0.219155  0.780845
      0.939594  0.0604063
      0.887168  0.112832
      0.589895  0.410105
      0.150417  0.849583
      0.673888  0.326112
      ⋮
      0.983417  0.0165835
      0.294901  0.705099
      0.994345  0.00565458
      0.858717  0.141283
      0.675784  0.324216
      0.0382285 0.961771
```

```
• # multinomial logit probabilities
• prob_actions = mapslices(softmax, u; dims=2)
```

```
cum_prob = 1000x2 Matrix{Float64}:
      0.947399  1.0
      0.219155  1.0
      0.939594  1.0
      0.887168  1.0
      0.589895  1.0
      0.150417  1.0
      0.673888  1.0
      ⋮
      0.983417  1.0
      0.294901  1.0
      0.994345  1.0
      0.858717  1.0
      0.675784  1.0
      0.0382285 1.0
```

```
• cum_prob = cumsum(prob_actions; dims=2)
```

```
• # will throw an error if we goof
• @assert all(cum_prob[:,2] . $\approx$  1)
```

```
y =
  Int64[0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1,      more ,
```

- *# instead of simulating random type-1 extreme values, we just*
- *# use a uniform variable and the CDF*
- `y = [searchsortedfirst(row, rand()) for row in eachrow(cum_prob)] .-1`

Dict{0 ⇒ 518, 1 ⇒ 482}

- `countmap(y)`

loglik (generic function with 1 method)

- `function loglik(y, X, theta)`
- `n,k = size(X)`
- `ff(z) = logcdf(Logistic(), z)`
- *# see footnote 6 on p. 778 in Greene 6th ed for this shortcut*
- `q = 2 .* y .- 1`
- `u1 = X*theta`
- `LL = sum(ff.(q.*u1))`
- *#*
- `return -LL # I *think* you'll need to flip sign to maximize`
- `end`

dloglik! (generic function with 1 method)

- *# note that the '!' means we're updating the first argument(s)*
- `function dloglik!(grad, y, X, theta)`
- *#*
- `n,k = size(X)`
- `u1 = X*theta`
- *#*
- *# create function*
- `ff(z) = cdf(Logistic(), z)`
- *#*
- *# all the broadcasting fuses operations into a single*
- *# loop instead of allocating temp vectors*
- *# this can help w/ speed + memory*
- `grad .= -vec(sum((y .- ff.(u1)) .* X; dims=1))`
- `return grad`
- `end`

dloglik (generic function with 1 method)

- *# wrapper to allocate gradient vector*
- `dloglik(y, X, theta) = dloglik!(similar(theta), y, X, theta)`

informationmatrix (generic function with 1 method)

- `function informationmatrix(y, X, theta)`
- *#*
- `n,k = size(X)`
- `infomatrix = zeros(k,k)`
- *#*
- `u1 = X*theta`
- `ff(z) = cdf(Logistic(), z)`
- `g = y .- ff.(u1) # as per Greene 6th ed p. 779`
- *#*
- `infomatrix = (g .* X)' * (g .* X)`
- *#*
- `return infomatrix # maybe flip signs?`
- `end`

g! (generic function with 1 method)

```
• # closures wrap likelihood & gradient
• begin
•   f(thet) = loglik(y,X,thet)
•   g!(grad,thet) = dloglik!(grad,y,X,thet)
• end
```

theta0 = Float64[0.0, 0.0, 0.0, 0.0]

```
• # initial guess
• theta0 = zeros(k)
```

Float64[-6.5473e-9, 1.80674e-8, -3.56951e-9, -1.54798e-8]

```
• # Check gradient against finite difference
• begin
•   fdgrad = finite_difference_gradient(f, theta0, Val{:central})
•   @assert fdgrad ≈ dloglik(y,X,theta0)
•   fdgrad .- dloglik(y,X,theta0)
• end
```

res = * Status: success

* Candidate solution
Final objective value: 3.916284e+02

* Found with
Algorithm: BFGS

* Convergence measures

$ x - x' $	$= 1.69e-08 \not\leq 0.0e+00$
$ x - x' / x' $	$= 8.09e-09 \not\leq 0.0e+00$
$ f(x) - f(x') $	$= 0.00e+00 \leq 0.0e+00$
$ f(x) - f(x') / f(x') $	$= 0.00e+00 \leq 0.0e+00$
$ g(x) $	$= 3.81e-10 \leq 1.0e-08$

* Work counters

Seconds run:	0 (vs limit Inf)
Iterations:	14
f(x) calls:	39
$\nabla f(x)$ calls:	39

```
• res = optimize(f, g!, theta0, BFGS(), Optim.Options(;show_trace=true))
```

	beta	betahat	tstat	se	pval
1	1.0	1.06057	10.043	0.105603	9.85822e-24
2	-2.0	-2.09389	-14.426	0.145147	3.54972e-47
3	1.0	1.08814	10.415	0.104479	2.11883e-25
4	0.5	0.586291	6.33893	0.0924905	2.31366e-10

```
• begin
•   theta1 = res.minimizer # should be about  $\beta$ 
•   vcov = informationmatrix(y, X, theta1)
```

```
•   vcovinv = inv(vcov)
•   stderr = sqrt(diag(vcovinv))
•   tstats = theta1 ./ stderr
•   pvals = map(z -> 2 .* cdf(Normal(), -abs(z)), tstats)
•
•   DataFrame(beta =  $\beta$ , betahat = theta1, tstat=tstats, se = stderr, pval=pvals)
• end
```
